


“Hello world”
of deep learning

Keras

If you want to learn theano:

http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/Theano%20DNN.ecm.mp4/index.html

[http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/RNN%20training%20\(v6\).ecm.mp4/index.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/RNN%20training%20(v6).ecm.mp4/index.html)



Very flexible

Need some
effort to learn

Interface of
TensorFlow or
Theano



keras

Easy to learn and use

(still have some flexibility)


You can modify it if you can write
TensorFlow or Theano

Keras

- François Chollet is the author of Keras.
 - He currently works for Google as a deep learning engineer and researcher.
- Keras means *horn* in Greek
- Documentation: <http://keras.io/>
- Example:
<https://github.com/fchollet/keras/tree/master/examples>

使用 Keras 心得


Deep Learning 研究生




朋友覺得我在



我媽覺得我在



大眾覺得我在



指導教授覺得我在




我以為我在



事實上我在

“Hello world”

- Handwriting Digit Recognition



MNIST Data: <http://yann.lecun.com/exdb/mnist/>

Keras provides data sets loading function: <http://keras.io/datasets/>

Keras


Step 1:
define a set
of function



Step 2:
goodness of
function



Step 3: pick
the best
function



(Because after the first layer, the layer's input dimension is equal to the pervious layer's output dimension.)

```
model = Sequential()
```

Only the first layer need to set the dimension of input.

```
model.add( Dense( input dim=28*28,  
                  output dim=500 ) )  
model.add( Activation( 'sigmoid' ) )
```

softplus, softsign, relu, tanh,
hard_sigmoid, linear

```
model.add( Dense( output dim=500 ) )  
model.add( Activation( 'sigmoid' ) )
```

```
model.add( Dense( output_dim=10 ) )  
model.add( Activation( 'softmax' ) )
```


Normally, the last layer's activation function is softmax.

Keras

Step 1:
define a set
of function

Step 2:
goodness of
function

Step 3: pick
the best
function



```
model.compile(loss='categorical_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])
```

Need to apply one-hot encoding on y.

Several alternatives: <https://keras.io/objectives/>

Keras



Set the method of how we update the function.

Step 3.1: Configuration

```
model.compile(loss='categorical_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])
```

SGD, RMSprop, Adagrad, Adadelata, Adam, Adamax, Nadam

Step 3.2: Find the optimal network parameters **Train the model**

```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

Training data
(Images)

Labels
(digits)

In the following slides

Keras

Step 1:
define a set
of function



Step 2:
goodness of
function




Step 3: pick
the best
function

Step 3.2: Find the optimal network parameters


```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

numpy array



Number of training examples

numpy array




Number of training examples

We do not really minimize total loss!


Mini-batch

Update the model per "batch" instead of per "all data".

Mini-batch



Mini-batch



- Randomly initialize network parameters

- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once

- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once

- ⋮
- Until all mini-batches have been picked

one epoch
Have processed all data

Repeat the above process


For each epoch, we will update the model (#examples/batch size) times.

Mini-batch

Batch size influences both *speed* and *performance*. You have to tune it.

```
model.fit(x_train, y_train, batch size=100, nb epoch=20)
```

Mini-batch



100 examples in a mini-batch

Batch size = 1 →
Stochastic gradient descent

- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
- ⋮
- Until all mini-batches have been picked


Repeat 20 times

one epoch

Two reasons that we shouldn't set the batch size too big:

1. The hardware limitations of GPU. (If the batch size is too big, the computation time for one batch will no longer stay the same.)
2. For DL, if we use full-batch, we will stuck at some local minima in the first few updates. (We need the randomness from mini-batch.)

Speed



∴ Error surface looks like this.

Very large batch size can yield worse performance

- Smaller batch size means more updates in one epoch

- E.g. 50000 examples

- batch size = 1, 50000 updates in one epoch

- batch size = 10, 5000 updates in one epoch

For one epoch

166s


In the same amount of time

1 epoch

17s


10 epoch

Process more data



Thanks to the parallel computation of GPU, we can spend the same time per update regardless of the batch size. (If batch size is not too big.) (Details on the next few pages.)

Speed - Matrix Operation




$$y = f(x) \quad \text{Forward pass (Backward pass is similar)}$$

$$= \sigma(W^L \dots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \dots + b^L)$$

Speed - Matrix Operation

- Why mini-batch is faster than stochastic gradient descent? **For stochastic gradient descent, we must wait for the previous update before calculating the next training example.**

Stochastic Gradient Descent



Mini-batch


Couldn't combine into one single operation.
matrix



Practically, which one is faster? **Mini-batch**

No matter how many elements in this matrix, we can use the same amount of time to solve it. (If it doesn't exceed the hardware limitation of GPU.)

Keras



Save and load models

<http://keras.io/getting-started/faq/#how-can-i-save-a-keras-model>

How to use the neural network (testing):

We have the label of the testing data: (predict + calculate the loss)

```
case 1: score = model.evaluate(x_test, y_test)
print('Total loss on Testing Set:', score[0])
print('Accuracy of Testing Set:', score[1])
```

We don't have the label of the testing data: (predict)

```
case 2: result = model.predict(x_test)
```

Keras

- Using GPU to speed training
 - Way 1
 - `THEANO_FLAGS=device=gpu0 python YourCode.py`
 - Way 2 (in your code)
 - `import os`
 - `os.environ["THEANO_FLAGS"] = "device=gpu0"`